

Moral Foundations of U.S. Political News Organizations

William E. Padfield¹, Erin M. Buchanan, Ph.D.², & Kayden N. Jordan, Ph.D.²

¹ Missouri State University

² Harrisburg University of Science and Technology

Author Note

William Padfield completed this project as part of his thesis requirements for a master's degree candidate in Psychology at Missouri State University. Submitted to Meta-Psychology. Participate in open peer review by commenting through [hypothes.is](https://tinyurl.com/mp-submissions) directly on this preprint. The full editorial process of all articles under review at Meta-Psychology can be found following this link: <https://tinyurl.com/mp-submissions>

Correspondence concerning this article should be addressed to William E. Padfield, 901 S. National Ave, Springfield, MO, 65897. E-mail: wpadfield@gmail.com

Abstract

Partisan differences and divisiveness have become an increasingly hot topic in psychological research. Many theories have been proposed to explain these differences and divisions including Moral Foundations Theory. The current research seeks to use a linguistic measure of Moral Foundations, the Moral Foundations Dictionary (MFD), to test the theory in terms of predicted partisan differences. In Study 1, we explore the MFD-Moral Foundations Questionnaire's measurement overlap using a multi-trait multi-method model. These latent trait models were used to inform possible improvement of the MFD. Next, we examined two replication extensions of previous work exploring the differences in political articles on moral differences. Through web scraping, we extracted articles from popular partisan news sources' websites, calculated MFD word frequencies, and identified words' respective valences. This process attempts to uncover news outlets' positive or negative endorsements of certain moral dimensions concomitant with a particular ideology. In Study 2, we gathered political articles from four sources. We were unable to reveal significant differences in moral endorsements, but we solidified the method to be employed in further research. In Study 3, we expanded their number of sources to 10 and analyzed articles that pertain to two specific topics: the 2018 confirmation hearings of U.S. Supreme Court Justice Brett Kavanaugh and the partial U.S. Government Shutdown of 2018-2019. Once again, no significant differences in moral endorsements were found. Together with past work, the results shed doubt on the validity of the MFD as a reliable measurement tool.

Keywords: politics, morality, psycholinguistics

Moral Foundations of U.S. Political News Organizations

The field of moral psychology has a long history with modern moral psychology beginning in the late 1960s with Lawrence Kohlberg's theory of moral development (Kohlberg & Hersh, 1977). Since then, Kohlberg's theory has been highly criticized with numerous theories and hypotheses proposed to replace it. One of these theories is Moral Foundations Theory (MFT) proposed by Haidt and colleagues (Haidt & Graham, 2007; Haidt & Joseph, 2004). While the theory itself has been widely criticized and debated, the goal of the current work is to tackle a related measurement issue, the Moral Foundations Dictionary (MFD, Graham, Haidt, & Nosek, 2009). Graham et al. (2009) developed the MFD - a lexicon-based linguistic measure of MFT - arguing that moral concerns and the propositions of MFT could be measured and tested with the words people use. Here we first review evidence and criticisms against this assumption. We then provide our own test of the MFD using partisan news articles.

Moral Foundations Measurement

Haidt and Joseph [Haidt, 2004] and Haidt and Graham (2007) formulated MFT as a method by which to capture the entirety of humans' moral domain. They argued that older theories of moral psychology were focused primarily on issues of justice, fairness, and caring, which are individually focused foundations of morality that align with the beliefs of political liberals. In other words, moral psychology ignored the valid moral foundations of conservatives. MFT holds that people's moral domain can be mapped by quantifying their endorsement of five moral foundations: *harm/care*, *fairness/reciprocity*, *ingroup/loyalty*, *authority/respect*, and *purity/sanctity* (Haidt & Graham, 2007). Thus, they proposed that the addition of these two additional domains could adequately capture the broader spectrum of morality.

Measurement of MFT has traditionally focused on two options: 1) the Moral Foundations Questionnaire (MFQ) - a typical Likert-type survey instrument that measures

moral relevance of each foundation to an individual and the moral judgment of various scenarios depicting each foundation, and 2) the Moral Foundations Dictionary (MFD) which was designed to capture language's role in moral reasoning and justification as used in speech and text (Graham et al., 2009, 2011a). The MFQ is a 30-item scale with 15 items on the moral relevance and judgment scenarios. The 15 moral relevance items are equally divided among the five foundations and examine how relevant a condition is to making a moral judgment on a scale of 1 (*not at all relevant*) to 6 (*extremely relevant*). These relevance items include examples such as: "Whether or not someone used violence (*harm*)," "Whether or not someone was denied his or her rights (*fairness*)," "Whether or not someone showed a lack of loyalty (*ingroup*)," "Whether or not an action caused chaos or disorder (*authority*)," and "Whether or not someone did something disgusting (*purity*)." The moral judgments items are also equally divided between the foundations and ask on a six-point scale how much one agrees with each of the statements. These judgment items include: "One of the worst things a person can do is hurt a defenseless animal (*harm*)," "Justice is the most important requirement of a society (*fairness*)," "I am proud of my country's history (*ingroup*)," "Men and women each have different roles to play in society (*authority*)," and "Chastity is an important and valuable virtue (*purity*)."

The internal consistency of the MFQ from Graham et al. (2011b) was $\alpha = .73$ averaged across subscales with a range of $\alpha = .65-.84$. Across six studies, the MFQ was found to have an average Cronbach's alpha of .63 for *harm*, .64 for *fairness*, .56 for *ingroup*, .59 for *authority*, and .71 for *purity* (Federico, Weber, Ergun, & Hunt, 2013; Graham et al., 2009; Graham, Nosek, & Haidt, 2012; Weber & Federico, 2013). Test-retest reliability was $r = .68-.82$ using a sample of 123 college students. Confirmatory factor analysis supported a well-fitted five-factor model (*harm/care*, *fairness/reciprocity*, *ingroup/loyalty*, *authority/respect*, and *purity/sanctity*) over two, individual (*harm* and *fairness*) versus group (*ingroup*, *authority*, and *purity*) foundations, or three, autonomy (*harm*, *fairness*), community (*ingroup*, *authority*), and divinity (*purity*) ethics, foundations factor model. The

five-factor structure also fit for non-Western samples, thus, providing evidence of the MFQ generalizability. Convergent validity was supported with correlations on other measures of morality (Graham et al., 2011b).

The MFD is composed of 259 words, with around 50 words assigned to each of the five foundations. Graham et al. (2009) created a preliminary list of words that they believed would be associated with the five foundations. Words such as *war* and *peace* should indicate a greater concern with *harm* foundation whereas words such as *homeland* and *terrorism* should indicate a greater concern with the *ingroup* foundation. The other foundation dictionaries include *equal* and *justice (fairness)*, *honor* and *protest (authority)*, and *holy* and *sin (purity)*. To validate the word sets, Graham et al. (2009) examined the frequency of MFD words in liberal and conservative sermons using the Linguistic Inquiry and Word Count [LIWC; Pennebaker, Booth, and Frances (2007)] computer program. They found liberal ministers used *harm*, *fairness*, and *ingroup* words more often than conservative ministers who used *authority* and *purity* words more often. Although conservative ministers were expected to use more *ingroup* words based on political ideology and previous research, an examination of the way liberal ministers used *ingroup* words revealed a tendency for the use of *ingroup* words to glorify rebellion and promote independence (i.e., the opposite direction from *ingroup* definitions). Effect sizes indicated relatively sizable difference between liberal and conservative sermons with Cohen's *d* values ranging from 0.56 to 1.27.

In contrast to this approach, Sagi and Deghani (2014) explored how moral words were used paired with other co-occurring concepts using Latent Semantic Analysis (LSA). They examined three different moral issues in different contexts to piece out specific moral words and their collocates. First, they looked at how moral words were used in relation to the World Trade Center compared to the Empire State Building in the New York Times from 1987-2007. After 9-11, the number of moral words associated with the World Trade Center increased, specifically *harm* words from the MFD. Second, they considered the changes in

how moral words were paired with mosque used in blogs as a response to the debate of building a mosque near Ground Zero following 9-11. They found words from the MFD were used more often with mosque during the main debate and then the co-occurrence decreased afterwards. Lastly, they examined moral language tied to the abortion debate in Congress. Republicans used more moral language overall; more specifically, Republicans tended to use more words associated with the *purity* foundation; while Democrats used more words associated with the *fairness* foundation.

Critiques of Moral Foundations

MFT has received criticism on the grounds that its assumptions regarding moral intuitions have little empirical basis. Suhler, Churchland, and Joseph (2011) criticized the content and taxonomy of the five foundations and question whether or not the foundations are sufficiently distinct as to stand as their own foundation. Likewise, Gray and Keeney (2015) argues that the measurement of the moral foundations confounds morality with other constructs (such as weirdness) and the moral foundations lack statistical and conceptual distinction. Schein and Gray (2015) argues that harm-based morality is the most parsimonious and logical explanation for MFT findings.

Beyond the critiques of the theory itself, many have also provided strong critiques of the measurement using MFD. Conducting two close conceptual replication studies using religious sermons and six extension studies using U.S. Congressional speeches, party platforms, State of the Union addresses, media outlets, and political organization texts, Frimer (2020) found that the liberal-conservative differences found in the original study replicated in only 30% of cases and the effect sizes were over 30 times smaller than reported in Graham et al. (2009). In a meta-analysis, the theorized differences were only found for authority and purity (e.g., conservatives used more authority and purity words). Loyalty was contrary to theoretical predictions (e.g., used more by liberals), and no differences were found for harm and fairness. Additionally, Frimer (2020) tested the theory under the most

forgiving conditions using not only the original MFD but also the revised MFD2 as well using log transformed scores and raw scores. Even under the best circumstances (log transformed MFD2 scores), only two of the foundations showed theorized differences but with very small effects. Other studies have also found conflicting findings which fail to support the basic assertions of MFT. Exploring the political Twitterverse, Sterling and Jost (2018) found differences in the MFD depend on political sophistication. For example, harm words were used more by liberals only if they had low levels of political sophistication; at high levels of political sophistication, conservatives used more harm words. While the MFD has not been extensively used (at least in published research), studies which have use the MFD to test partisan differences have found mixed results in terms of supporting MFT (Clifford & Jerit, 2013; Frimer, Tell, & Haidt, 2015; Sagi & Dehghani, 2014). Like Frimer (2020), we attempt to give the MFD the best chance for success by incorporating a weighting method to address the skewedness and sparsity of the linguistic data.

The most successful attempts to validate a linguistic measure of moral foundations have used more complex methods (compared to the original dictionary-based approach). These attempts generally fall into two categories: human annotations and semantic representations. Hopp, Fisher, Cornell, Huskey, and Weber (2021) represents the first approach where instead of relying on ‘experts’, a large sample of lay people were crowdsourced to manually annotate document for content relevant to each foundation. While resource-intensive, this method better captures how people may encounter and judge moral issues in everyday life leading to a more ecologically valid dictionary. Garten, Boghrati, Hoover, Johnson, and Dehghani (2016) represents the second categories which instead relies on data-driven semantic analysis. Using shallow neural net models like word2vec (Mikolov, Chen, Corrado, & Dean, 2013), moral foundations were measured using semantic similarity vectors rather than simple word counts. Here, we try to augment the simpler approach to add more context to see if the MFD liberal-conservative differences can be replicated.

First, we use the MFQ and MFD in a multi-trait multi-method approach to determine if the dictionary can be improved and mapped into the theoretical MFT traits. This study examined the measurement overlap between these two sources, as well as the latent traits after controlling for the overlap. Using the results from this study, we examined two conceptual replications of Frimer (2020). We test the liberal-conservative difference proposed by MFT using partisan media content and the MFD. We first test differences in a general news corpus compared liberal sources (National Public Radio [NPR] and the New York Times [NYT]) to conservative sources (Fox News and Breitbart). Second, we examined news about specific political events: Brett Kavanaugh’s U.S. Supreme Court nomination and the U.S. Government Shutdown in 2018-2019. In the second study, the news sources were expanded to cover NYT, NPR, Slate, Huffington Post, and Politico on the liberal side and Fox News, Breitbart, The Rush Limbaugh Show, The Blaze, and Sean Hannity on the conservative side. To address potential issues of measurement with the MFD, the MFD was combined with the information produced in Study 1 and weighted by valence to create weighted percentages to better specify endorsement. Like Frimer (2020), results were judged based on the direction of the liberal-conservative difference compared MFT predictions and the size of the effect compared to past studies.

Study 1

We approached this study with the intention to answer a method question. That is, this portion of the current research was conducted in order to solidify the best method by which to analyze political news text under the MFT framework while also alleviating some of the aforementioned potential dictionary issues with the MFD. In this set of studies, we examined the MFD and MFQ within the same study. The first two sub-studies were used to develop a (potentially) larger dictionary for the MFD and collect data on both scales. The original dictionary was developed conceptually, identifying words that *should* go with each moral trait. In Study 1.1 and 1.2, we collect data to expand the dictionary. The last

sub-study examined the MFD and MFQ using multi-trait multi-method (MTMM) analysis using guidelines from Byrne (2001) and Widaman (1985).

Study 1.1

In this study, participants were given a word association task elicit additional semantically related word forms for each moral foundation. This data was used to expand the MFD for each latent area to determine if these additional words could improve scores and test measurement relation to the MFQ in Study 1.3.

Participants. 466 participants were assessed from a large Midwestern university. Participants were given course credit for their introductory psychology course for completing the study. 14 participants had missing data on the MFQ and were excluded from analyses. The MFQ data was screened for multivariate outliers with Mahalanobis distance as described in Tabachnick and Fidell (2012), and 15 outliers were found using $\chi^2_{p<.001}(15) = 37.70$. These participants were excluded from further analyses, representing 437 final participants. Participants were asked to denote their political party, and 25.9% indicated they were Democrats, 42.8% were Republican, and 31.1% indicated they were Independent. 0.2% of values were missing for this question.

Materials and Procedure. A complete example of the survey can be found online at <https://osf.io/6mv2g>. First, participants were given a description of associative memory as the relation between words that comes about through many pairings in writing and speech. Next, the free association task, similar to that used in Nelson, McEvoy, and Schreiber (2004) and De Deyne, Navarro, Perfors, Brysbaert, and Storms (2019) was described to the participants as listing the “first word that pops into mind”. The participants were then given three example free association cues, *lost*, *old*, and *article*. For each cue, participants were asked to write all the words that come to mind. To elicit free association to the moral foundation areas, participants were given the following instructions:

“Moral Foundations Theory states that when making moral judgments/decisions, the

concerns people have can be divided into five categories. Below are labels of each of these five categories. You will then be asked to list words you think are associated with each of the labels.”

Each of the foundation pairs were listed together (i.e., *harm/care*, *fairness/reciprocity*, etc.) with a space for participants to list their free association concepts. After the free association task, participants were then given the 15-items from the moral relevance section of the MFQ as described in the introduction. Last, participants were asked to denote their political orientation from 1 *conservative* to 10 *liberal*, as well as which political party they associated with: Democrat, Republican, and Independent. The survey was delivered through Qualtrics, and participants were recruited through the online participant management system for the university (SONA). Each participant signed an online consent form at the beginning of the study and was given participation credit at the end of the study.

Results. All data was screening for inaccurate responses, as well as missing data, as described in the participant section. The final MFQ data was screened for assumptions of normality, linearity, and homoscedasticity. The MFQ data was slightly skewed but otherwise linear and homoscedastic after accounting for skew. No transform was applied. The sum of each moral foundation area was calculated in order to determine which words were linked to their respective moral foundation. The average scores were: harm ($M = 14.18$, $SD = 2.44$), fairness ($M = 14.30$, $SD = 2.59$), ingroup ($M = 12.74$, $SD = 2.92$), authority ($M = 12.05$, $SD = 3.01$), and purity ($M = 11.57$, $SD = 3.36$). Participants free association responses were processed using the *tm* (Feinerer, Hornik, & Meyer, 2008) and *corpus* (Perry, 2024) libraries after manual correction of spelling errors. Each set of answers was cleaned for punctuation, English stop words (e.g., *the*, *an*, *of*) were removed, and each word was stemmed using the English library in *tm*. We did not combine related words in this section (i.e., *injure* and *injury*, which have different stems *injur* and *injuri*) to allow for maximum coverage of different word forms present in the dictionary. Additionally, with the use of automated stemmers like that present in the *tm* library, leaving both word forms in the

dictionary would capture more of the concepts present in future analyses with a different corpus without the requirement on the coder to manually recode all word forms. Frequency counts of the stemmed words were tabulated and only words mentioned with at least one percent frequency were used in the subsequent analyses. The complete set of word frequencies for each foundation can be found in our supplemental materials.

For clarity, the processing pipeline for each separate domain was:

1. Manual spell checking of free association responses.
2. Word tokenization and stemming.
3. Create a term by document matrix.
4. Remove punctuation and stopwords via the *tm* stopwords list.
5. Calculate percent frequency across documents.
6. Remove words with less than 1% mentions across participants.

This procedure generated a large frequency of words for a new dictionary of moral foundations: harm 96, fairness 77, ingroup 87, authority 84, and purity 81. These concepts were included in the full dictionary used for Study 1.3. We additionally created a reduced dictionary that included only concepts correlated with their respective moral foundations scores. Correlations between word frequency and the sum of the MFQ were calculated for each foundation and set of concepts. Words were included in the reduced dictionary if their correlation was two standard deviations away from the mean correlation for that foundation. The reduced dataset included the following number of words for each foundation: harm 4, fairness 3, ingroup 2, authority 4, and purity 1.

Study 1.2

In this study, participants were asked to write about chemical warfare (wave 1), abortion, same-sex marriage, and environmentalism (wave 2). These prompts were designed to elicit writing about moral concepts that would relate to the moral foundations areas for

expansion of the dictionary.

Participants. Participants were recruited in two waves as part of a larger investigation on priming political and religious attitudes. Participants were recruited via an online research system (SONA) and were given course credit for their participation. 463 participants were included in the this study. The study was mostly women (53.9%) and White (76.4%) participants with a mix of minority participants: Black (6.1%), (3.6%), (4.2%), Native American (1.9%), Mixed (2.9%) and Other (4.9%). The average listed age was 19.75 ($SD = 2.94$).

Materials and Procedure.

Wave 1. Data was again collected via Qualtrics. Four fake new stories were presented to participants, which were roughly 400 words each. First, all news stories included a few sentences describing the use of chemical weapons in the Syrian civil war. The news stories were manipulated with political (Republican v. Democrat) and religious (religious v. not) quotes in a 2 x 2 between-subjects design. News stories can be found in the online materials. After consenting to participate in the study, participants were randomly shown one of the four new articles about Syria's use of chemical weapons. Participants were then asked to write for 5-10 minutes about their reaction to Syria's use of chemical weapons and the needed response from the United States. Participants also completed the 30-item version of the MFQ as described previously. In addition to basic demographics (gender, age), participant political orientation was assessed with the same scale described in Study 1.1.

Wave 2. The second wave of data collection included different writing prompts designed to capture more of the moral foundation areas in their writing. The following writing prompt(s) were used, "Please write about your attitudes on abortion (or same-sex marriage or environmentalism) as well as your reason for this stance." The three prompts were chosen to create a more varied word set by using topics that should elicit words from each moral foundations category by soliciting a moral response. The writing prompts were randomized within participant, and each participant completed all three prompts.

Participants then completed the MFQ, demographics, and the political orientation scale.

Results. Participant data was first spell checked and screening for inaccurate responses. Participants who did not write more than fifty words in response to a given prompt were excluded ($n = 51$). One missing datapoint was estimated using the *mice* library from *R* for a missing MFQ question (Buuren & Groothuis-Oudshoorn, 2011), and all other missing data was present in the demographics sections, which were not filled in. The MFQ data were screened for outliers using Mahalanobis distance, and 25 outliers were found at $\chi^2_{p<.001}(44) = 78.75$. These data were excluded leading to a final sample size of 387. Data were screened for assumptions described in Study 1.1 and were found to be satisfactory.

In the first study, only free association responses were collected, but in this study, full writing samples were collected. Therefore, we expected many of the words listed to be part of creating a cohesive discourse, rather than only related to the moral foundation targeted. To find only the most related words, the correlation between word frequency and MFQ subscore was calculated, and words with correlations greater than two standard deviations outside the mean were selected for the dictionary analysis in Study 1.3. The sum of each MFQ subscore was calculated for the correlation analysis. The average scores were: harm ($M = 13.85$, $SD = 2.79$), fairness ($M = 14.12$, $SD = 2.78$), ingroup ($M = 11.96$, $SD = 3.24$), authority ($M = 11.52$, $SD = 3.06$), and purity ($M = 11.69$, $SD = 3.61$).

The correlated set of words was then coded by two of the authors (EB, WP) for relevancy to the MFQ subscore and all disagreements were discussed to determine final inclusion. For example, for harm words in this study both *also* and *Syria* appeared as correlated words to the harm foundation. The first word does not meet face validity of correlation with the dictionary, and these type of function words or ambiguous verbs (i.e., *get*, *can*) were usually excluded. The second example is likely because of the writing prompt, and therefore, proper nouns were also excluded. The reduced dataset included the following number of words for each foundation: harm 11, fairness 11, ingroup 12, authority 17, and

purity 29.

Study 1.3

MTMM analyses use latent variable modeling to represent the underlying structure of a proposed phenomenon, while also accounting for the measurement of that phenomenon. As shown in Figure 1, MTMM models are programmed to represent the traits that are measured (left hand side circles) and the methods used to measure those traits (right hand side circles). The squares in the middle represent the actual observed data, which is the MFQ questions and MFD scores. In the first model, *correlated traits correlated methods*, the latent variables are allowed to correlate on the traits and methods side, but not cross correlate from traits to methods. All models are then compared to this model to examine fit, convergent, and divergent validity. In the second model, *no traits correlated methods*, the trait side of the model is deleted and only methods are considered. The model should significantly degrade in fit, as the traits should be useful in explaining model variance for convergent validity (i.e., you do not want a model that suggests all the variances is due to measurement).

In the third model, *perfectly correlated traits freely correlated methods*, the traits are added back to the model, but each correlation is set to one. The fit is compared to model 1, indicating if divergent validity of the traits has occurred. A significant degrade in fit should be found to indicate that trait correlations are not perfect. Last, model four *correlated traits uncorrelated methods*, the traits are allowed to naturally correlate but the methods correlations are set to zero. In comparison with model 1, the model should be roughly equal for divergent validity. This result would indicate that the methods each represent a different component to the measurement of the model, and do not represent the same measurement twice.

Given the goal of improved measurement, we compared model 1 of the MTMM steps across several dictionaries to determine the best measurement for the dictionary. To assess

model fit, we used the Root Mean Squared Error of Approximation (RMSEA, Steiger, 1990), the Root Mean Squared Residual (RMSR, Jöreskog, 1971), the Confirmatory Fit Index (CFI, Bentler, 1990), the Tucker-Lewis Index (TLI, Bentler & Bonett, 1980). The RMSEA and RMSR should be low for adequate model fit ($< .10$) and the CFI and TLI should be high for adequate model fit ($> .90$). For model comparison, we used a Δ CFI of greater than .01 to indicate a significant degradation in fit (Hu & Bentler, 1999). We will additionally use the standardized loadings for each of the traits to determine which model has the best measurement of moral foundations ($> .30$ desired).

Method.

MFD Dictionaries. We created three versions of the MFD for separate MTMM analyses. The first dataset *original MFD* was comprised of the original concepts from the MFD. We added all versions of the original words (i.e., abuse \rightarrow abusive, abuser, abused) to ensure all forms of the words were captured within the stemming procedure. These words were then stemmed using the same procedures as the data processing described for the overall study to match the data processing completed on the participant prompt responses in Study 1.2. The *complete MFD* dataset included the original MFD concepts, all concepts listed from Study 1.1, and concepts from Study 1.2. The *reduced MFD* dataset included all concepts from the original MFD, correlated concepts from Study 1.2, and concepts from Study 1.2 (only correlated concepts were selected in Study 1.2). The percent of concepts within each moral found area was calculated on the writing prompts from Study 1.2. The percent values were calculated on the processed data excluding all stop words. These values are represented as the MFD squares in Figure 1.

MFQ Subscore. The MFQ individual questions for each participant were used from Study 1.2. These values are represented as the MFQ Qs squares in Figure 1. The final number of participants included in this model was 387.

Models. For clarity, we programmed the following models:

1. Model 1 *Correlated Traits Correlated Methods*

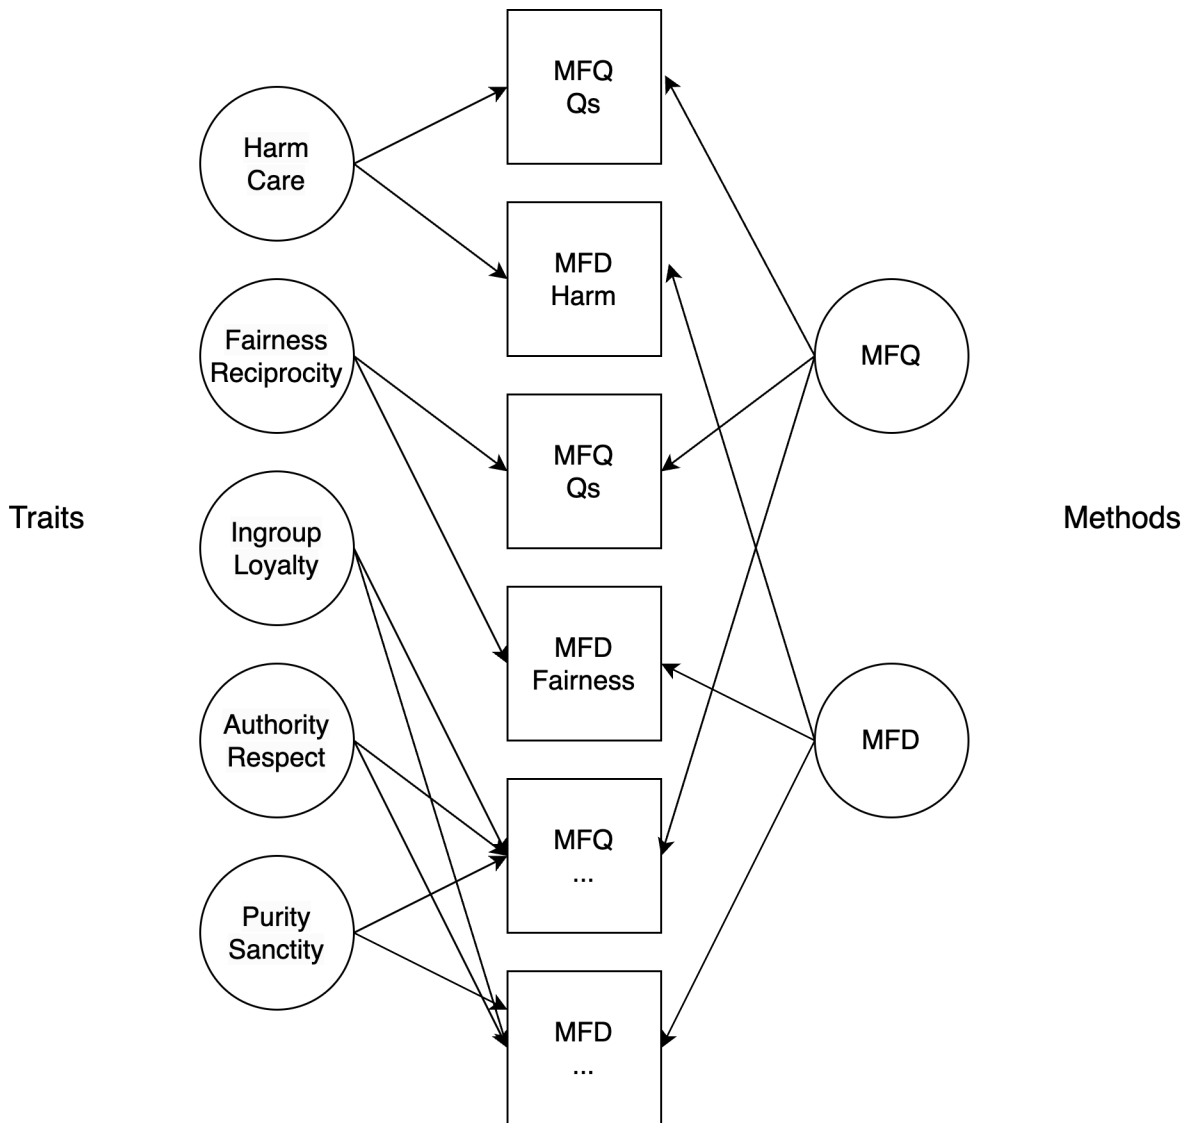


Figure 1. Representation of the multi-trait multi-method model for the MFD and MFQ. Not all measured variables (squares) are shown for ease of reading. The left hand side represents the traits or foundation areas of MFT, while the right hand side represents the two measurement tools for the MFT.

- a. *Original* Original MFD Words
 - b. *Complete* Original MFD Words + Study 1.1 + Study 1.2 Words
 - c. *Reduced* Original MFD Words + Study 1.1 Reduced + Study 1.2 Words
 - d. Determine if model suggests a feasible structure (fit indices adequate + loadings adequate)
 - e. Pick the best model to continue (using model comparison)
2. Model 2 *No Traits Correlated Methods*
 3. Model 3 *Perfectly Correlated Traits Correlated Methods*
 4. Model 4 *Correlated Traits Uncorrelated Methods*

Results. Table 1 indicates the average percents found for each dictionary within each moral foundation area. The original MFD shows the lowest percentage of moral words for each foundation area, while the complete and reduced dataset show increases in the percent of text that would be considered moral words.

Model 1. Fit indices from the MTMM Model 1 step are shown in Table 2. The complete parameter estimates from these models are saved online for full inspection due to their length. The fit indices from these models indicate that all models showed good fit. The original and complete MFD models showed equal fit ($\Delta \text{CFA} \leq .01$), and both models were better than the reduced MFD model. We then inspected the loadings to determine the adequacy of the loadings for the MFQ and MFD. These results are shown in Table 3. All but one question for the MFQ load adequately ($>.30$) on each of their moral foundation areas. None of the MFD areas load onto their moral foundation area, as all loadings are close to zero (and non-significant, see online). The excellent model fit was likely due to the structure of the MFQ, as the MFD results do not show relation to their foundation area. We did not run further models, as the goal was to show that Model 1 was the best model with appropriate loadings for foundations.

Discussion

Study 1 was designed to investigate if the original MFD could be improved with additional concepts. As shown in Table 1, the original MFD is often heavily skewed toward zero with very low percentages of moral words used in text. These results were found even when writing about concepts that should theoretically activate moral values. The suggested improvements did increase percentages of included words; however, the results from the MTMM indicate that the MFD does not appear to relate to the MFQ moral foundations. In a separate set of studies, we examined if the original MFD proposed results for liberal and conservative values could be found within new sources (Sagi & Dehghani, 2014). Further, we examined if the MFD could be improved with markers of valence (i.e., positive, negative), rather than just inclusion of more words.

Study 2

Method

We hypothesized the news sources generally perceived as liberal leaning (*NPR* and *The New York Times*) would contain MFD words and valences indicating endorsements of the individualizing moral foundations (*harm/care* and *fairness/reciprocity*). Additionally, we hypothesized the two sources generally perceived to be conservative leaning (*Fox News* and *Breitbart*) would feature MFD words and valences indicating equal endorsement of all five foundations.

Sources. Political articles were collected from the websites of four notable U.S. news sources, a process known as web scraping. The sources were *NYT*, *NPR*, *Fox News*, and *Breitbart*. They were selected for their widespread recognition and the fact political partisans have strong preferences for some sources over others. We determined the political lean of each source by referencing Mitchell, Matsa, Gottfried, and Kiley (2014)’s article demonstrating the self-reported ideological consistency represented by the consumers of several news sources. In general, the *NYT* and *NPR* are preferred by consumers reporting a

liberal bias or lean. In contrast, *Fox News* and *Breitbart* are believed to have a conservative bias or lean. Mitchell et al. (2014)’s article presented political ideology as a scale ranging from “consistently liberal” to “consistently conservative.” In between these extremes lie more moderate positions, including “mostly liberal,” “mixed,” and “mostly conservative.” Owing to the lower number of sources analyzed herein, we elected to categorize the sources as either “liberal” and “conservative” in order to form a basis for comparison. At the time of our data collection, these sources were also selected because they were open (i.e., no subscription required to access articles), and their websites were designed in a way that made webscraping possible.

Political articles in particular were identified and subsequently scraped by including the specific URL directing to each source’s political content in the *R* script. For example, rather than scrape from *nytimes.com*, which would return undesired results (non-political features, reviews, etc.), we instead included *nytimes.com/section/politics* so that more or less exclusively political content was obtained. All code for this manuscript can be found at <https://osf.io/5kpj7/>, and the scripts are provided inline with this manuscript written with the *papaja* library (Aust & Barth, 2017).

Identification of the sources’ political URLs presented a problem for two of the sources owing to complications with how their particular sites were structured. While in the multi-week process of scraping articles, we noticed word counts for *NPR* and *Fox News* were not growing at a similar pace as those from the *NYT* and *Breitbart*. Upon investigation, we found another, more robust URL for political content from NPR: their politics content “archive.” The page structure on NPR’s website was such that only a limited selection of articles is displayed to the user at a given time. Scraping both the archive and the normal politics page ensured we were obtaining most (if not all) new articles as they were published. We later ran a process in order to exclude any duplicate articles. *Fox News* presented a similar issue. We discovered *Fox News* utilized six URLs in addition to the regular politics

page. These URLs led to pages containing content pertaining the U.S. Executive Branch, Senate, House of Representatives, Judicial Branch, foreign policy, and elections. Once again, duplicates were subsequently eliminated from any analyses.

Materials. Using the *rvest* library in the statistical package *R*, we pulled body text for individual articles from each of the aforementioned sources (identified using CSS language) and compiled them into a dataset (Wickham, 2016). Using this dataset, we identified word count and average word count per source. This process was completed once daily starting in February 2018 until March 2018. Starting in mid-March 2018, the process was completed twice daily - once in the morning and again in the evening. Data collection was terminated once 250,000 words per source was collected in April 2018.

Data analysis. Once data collection ended, the text was scanned using the *ngram* package in *R* (Schmidt, Gonzalez-Cabrera, & Tomasello, 2017). This package includes a word count function, which was used to remove articles that came through as blank text, as well as to eliminate text picked up from the Disqus commenting system used by certain websites. At this point, duplicate articles were discarded.

The article text was processed using the *tm* and *ngram* packages in *R* in order to render the text in lowercase, remove punctuation, and fix spacing issues (Feinerer & Hornik, 2017). The individual words were then reduced to their stems (i.e., *abused* was stemmed to *abus*). The same procedure was applied to the MFD words and the words in the Warriner, Kuperman, and Brysbaert (2013) dataset. Using the Warriner et al. (2013) dictionary, the words making up each of the five foundations in the MFD were matched to their respective valence value.

The Warriner et al. (2013) data includes nearly 14,000 English lemmas that have been rated for their valence, arousal, and dominance (mirroring Bradley & Lang, 1999). Word emotion ratings have been used to estimate the sentiment of a text - a very popular task within classification research - by generally averaging the valence scores of the the words

matched between the text and the available norms (Leveau, Jhean-Larose, Denhière, & Nguyen, 2012). Other suggestions include the simple summation of word's positive or negative valence from a text or respective weighting based on word frequency (Hutto, 2021; Loria, 2020). In each of these cases, an overall polarity is desired, which is the case here, creating an overall score for each of the five MFD domains. Sentiment coding packages often score items from -1 to 1 or other values around zero wherein negative scores represent negative sentiment, while positive scores represent positive sentiment. The Warriner et al. (2013) dataset ranges from 1 (*unhappy*) to 9 (*happy*). Thus, to ensure that we anchored their valences around zero, we *z*-scored the dataset so that negative scores represented values below the average sentiment and positive scores represented values above the average sentiment.

Given the results from Study 1.3, we used the original MFD to calculate the number of words included for each moral foundation area. We could also suggest using the larger complete dataset, given that the CFI indicated similar models, and we note here that the pattern of results do not change if that dictionary is used instead. The total number of times each of those words was used in an article was calculated. This information was merged with the valence values from the Warriner et al. (2013) dataset. A weighted score was calculated by multiplying the frequency of occurrence within a document times the *z*-scored valence values and divided by the word count for that article. The final score was multiplied by 100 to represented a weighted percentage. Finally, the weighted scores were summed within in each article to get a total weighted percentage of the use of moral area specific words by article. Valences were *z*-scored in order to eliminate any ambiguity regarding the direction of the valence. Positive values indicate positive valence, and negative values indicate negative valence. Words were categorized in accordance to their MFD affiliation, creating a weighted sum for each moral foundation. See Table 4 for an example.

Results

To analyze if news sources adhered to differences in word use based on their target audience, we utilized a multilevel model (MLM) to analyze the data. MLM is a regression technique that allows one to control for the repeated measurement and nested structured of the data, which creates correlated error (Gelman, 2006). Using the *nlme* library in *R* (Pinheiro, Bates, Debroy, Sarkar, & Team, 2017), each foundation's weighted percentage was predicated here by the political lean of the news source, using the individual news sources as a random intercept to control for the structure of the data. Therefore, the models were calculated by: $\text{Weighted Score} \sim \text{Political Lean} + (\sim 1 | \text{Source})$.

The multilevel model did not indicate the presence of any significant or practical effect of political lean for any of the five moral foundations. The strongest effect size was observed for the *authority/respect* foundation, but the effect was in the opposite direction from what was originally hypothesized - liberal sources tended to use more *authority/respect* words than did conservative sources. Descriptive and test statistics, *p*-values and effect sizes (Cohen's *d*) can be found in Table 5. To interpret the weighted scores, one can examine the mean and standard deviations for each. A zero score for the mean, with a non-zero standard deviation, would indicate a perfect balance of positive and negative words in each category, likely representing a neutral tone when all words are considered. Negative percentages would indicate more representation of the negative words in the MFD area, while positive percentages indicate an endorsement of the positive words in a MFD. Therefore, we suggest using the sign of the mean score to determine the directionality of the endorsement for the MFD (positive, neutral, negative), and the standard deviation to ensure that a zero score is not zero endorsement (i.e., a SD of zero indicates no words were used). Based on the weighted percent values for the five foundations, we observed that MFD words seem to make up a small portion of the article text. Furthermore, the observed percentages and means appear to indicate a generally positive endorsement of all five foundations across both liberal

and conservative sources.

Discussion

The results obtained in Study 2 fail to support the differences predicted by MFT. First, differences between liberal and conservative news sources failed to reach statistical or practical (effect size) significance for any foundations. Second, looking purely at the direction of the differences, the results run completely contrary to the differences expected by MFT. Conservative sources scored higher on concern for harm and fairness while liberal sources scored higher on concern for loyalty, authority, and purity. Overall, the results lend strong support to the case against the MFD and weakens the case generally for MFT.

However, one possible limitation of the current data is the generalness of the corpus. The selection of the broad and amorphous topic of “political news” may have led to the scraping of large numbers of articles with little to no moral-centric content. To address this limitation, two changes that were subsequently employed in Study 3. First, we elected to include more news sources for web scraping and analysis in addition to the four used in Study 2. Second, we chose to focus data collection efforts on two heavily moralized events in the Trump administration: (1) the nomination and confirmation of Justice Brett Kavanaugh to the U.S. Supreme Court and the U.S. government shutdown in December 2018 through January 2019 over disagreements about funding a U.S.-Mexico border wall. Hence, Study 3 is the best test of the MFT and the MFD in the context of partisan news.

Study 3

Method

In contrast to Study 2, we approached Study 3 with the intention to confirm the method employed was valid for the analysis of the scraped text as well as for any inferences drawn from the analyses. For the second study, we hypothesized that news sources perceived as liberal will exhibit positive endorsements of the individualizing moral foundations

(*harm/care* and *fairness/reciprocity*) in their articles reporting on both the Kavanaugh confirmation hearing as well as the 2018-2019 government shutdown. News sources perceived as conservative are hypothesized to positively endorse all five foundations equally in their coverage of the Kavanaugh hearing and the government shutdown. We tested the hypothesis by analyzing the content scraped from news sources' web pages spanning the two weeks before (September 13, 2018) and two weeks after (October 11, 2018) Kavanaugh's confirmation hearing, owing to its prominence in the news. Likewise, we analyzed content spanning two weeks before the start of the government shutdown (December 8, 2018) to two weeks following the end of the shutdown (February 8, 2019). The content will be analyzed for valence and moral alignment under MFT.

Sources. Articles pertaining to the Brett Kavanaugh Supreme Court confirmation hearing and the 2018-2019 U.S. Government shutdown were scraped from the websites of 10 U.S. news sources. As in Study 2, these sources were selected owing to their favorability among political partisans according to Mitchell et al. (2014). The sources favored by the highest proportion of consistent liberals were the *NYT*, *NPR*, *Slate*, *Huffington Post*, and *Politico* (Mitchell et al., 2014). The sources favored by the highest proportion of consistent conservatives included *Fox News*, *Breitbart*, *The Rush Limbaugh Show*, *The Blaze*, and *Sean Hannity*. These sources were primarily selected due to their political lean but also due to our ability to webscrape their contents without a subscription based service. Political articles referencing Brett Kavanaugh's nomination process were identified and subsequently scraped by including the URL for each source's coverage of the nomination in the *R* script.

Materials. Using the *rvest* and *RSelenium* libraries, we pulled body text for individual articles from each of the aforementioned 10 news sources and compiled them together (Harrison & Kim, 2020; Wickham, 2016). Using this dataset, we identified word count and average word count per source. This process was run for articles pertaining to Kavanaugh's nomination that were published between September 13, 2018 and October 11, 2018 inclusive. This date range was selected in reference to the widely-publicized and viewed

nomination hearing on September 27, 2018. We set the start date at September 13 (two weeks before the hearing) and the end date at October 11 (two weeks after the hearing) so that we could capture a large amount of data (roughly one month) during which Kavanaugh's nomination was at its peak saturation in news coverage.

The same process was followed for scraping articles related to the partial U.S. Government shutdown of 2018-2019. The articles scraped were published between December 8, 2018 and February 8, 2019 inclusive. Once again, we elected to scrape articles published two weeks before and after the event in question in order to capitalize on the shutdown's saturation in American news media.

Data analysis. As in Study 1 and 2, the text was scanned with *ngram* (Schmidt et al., 2017). Again, blank articles, text from the Disqus system, and duplicate articles were removed. The text was processed and stemmed in order to convert to a usable form for further analysis (Feinerer & Hornik, 2017). Each individual word was reduced to its stem (i.e., *diseased* was stemmed to *diseas*). Once again, the same procedure was applied to the MFD words and the words in the Warriner et al. (2013) dataset. We obtained the words' percent occurrence in the text. Percents were multiplied by *z*-scored valence and categorized into their proper MFD category.

Results

We utilized a multilevel model (MLM) to analyze if news sources leveraged different vocabularies based on target audience in the same manner as Study 2. Each foundation's weighted percentage was predicted by the source's political lean, using the individual source as a random intercept to control for the nested data structure. Two separate MLMs were constructed from datasets compiled for each topic of interest: the Kavanaugh hearing and the partial government shutdown of 2018-2019.

For the Kavanaugh topic, the multilevel model indicated the presence of a statistically

significant effect for *harm/care*, but the practical effect denoted by Cohen's d was found to be small. There were no other significant or practical effects of political lean for any of the other four moral foundations. The effect for *harm/care* was in the hypothesized direction with liberal sources tending to use more positively *harm/care* words than conservative sources. Descriptive and test statistics, p -values and effect sizes (Cohen's d) can be found in Table 6.

For news articles about the partial U.S. Federal Government Shutdown of 2018-2019, there were no significant or practical effects of political lean for the moral foundations (see Table 7). A small-to-medium effect size was observed for *authority/respect*. The effect was in the predicted direction as well, as conservative sources tended to offer more positive endorsements of the foundation than liberal sources. This effect was similar to Study 2 in which the largest effect size was observed for *authority/respect*. As noted before, the effect found in Study 2 was in the opposite direction as what was hypothesized. Thus, it is difficult to draw comparisons between the two studies despite the similar pattern for effect size. Owing to a lack of similar effects for either Study 2 or the Kavanaugh topic, there is doubt as to whether or not a practical or generalized effect exists for *authority/respect*.

Based on the weighted percent values for the five foundations applied to both topics, MFD words seem to make up little of the article text. A similar pattern was observed for the results in Study 2. As in Study 2, the percentages and means seem to indicate a generally positive endorsement of all five moral foundations across both political leanings.

Discussion

The results obtained in Study 3 are slightly more in line with expectations from MFT, however, in all but one case, the differences were not statistically significant. The only statistically significant difference was in the harm foundation; liberals did endorse harm more than conservatives during the Kavanaugh hearings, but the effect was smaller than expected given Graham et al. (2009). Generally for both the Kavanaugh hearings and the Government

shutdown, the differences were in the expected direction (liberals endorsing harm more and conservatives endorsing authority and loyalty more), but the differences were not statistically significant and the practical effects much smaller than reported in Graham et al. (2009).

While not significant, fairness was endorsed more by conservatives in both cases. In contrast, purity was endorsed more by conservatives for the government shutdown, but more by liberals for the Kavanaugh hearing though neither difference was statistically significant.

Together with Study 2 and past work, Study 3 further calls into question the usefulness and efficacy of the Moral Foundations as well as cast doubt on the predictions made by MFT. One issue that Study 3 brings up is that moral differences between liberals and conservatives may be moderated by a number of factors as shown in Sterling and Jost (2018). For example, here we have liberals concerned with purity for the Kavanaugh hearings (a moral bad for liberals and a moral good for conservatives) and conservatives concerned with purity during the government shutdown (a moral bad for conservatives and a moral good for liberals). The original predictions made by the Moral Foundations are very simplistic and likely much more complex in reality.

Conclusions

Within the theoretical framework of MFT (Haidt & Graham, 2007), we attempted to improve the MFD (Graham et al., 2009) and demonstrate its relation with other reliable measurements of moral foundations, the MFQ. In Study 1, we examined potential new concepts to add to the MFD by collecting word associations and moral related texts from participants. Several dictionary improvements were suggested, but the results of a multi-trait multi-method analysis indicate that the MFD does not relate to the moral foundations found in the MFQ. Therefore, we did not find support for the use of the MFD in relation to the MFT using a convergent and divergent validity technique.

In Study 2 and 3, we attempted to devise a method leveraging the MFD in order to

quantify political bias stemming from content published by several prominent American news sources. In Study 2, the results obtained were not significant in any statistical or practical sense. There was a small effect found for the *authority/respect* foundation according to Cohen's d , but the effect suggested endorsement for that foundation opposite to what was hypothesized. In Study 3, the direction of the differences were generally in line with MFT, however, the difference were very small and lacks statistically significance. Building on Frimer (2020) and other critiques of MFT and the MFD, we show that MFD is not a useful tool for measuring moral language while also calling into question the validity of MFT in terms of theorized partisan differences.

Despite the fact the results regarding political bias were inconclusive, we still retain confidence in the overall structure of the methodology established in the current study. Specifically, the procedure for scraping text from the web, processing, stemming, and weighting the scores with valence seems to represent a solid method for preparing a high quantity of text passages for data analysis. We implemented valence as an indicator of the directionality of endorsement due to the inherent ambiguity of simply calculating MFD word percent occurrence. This method served both to augment the face validity of the MFD by incorporating valence (thus reducing ambiguity) and to generate a score that is easy to understand and analyze.

However, political lean and other important political/moral constructs may be communicated through methods other than through the use of particular words. As mentioned before, MFD is a measure developed from MFT (Graham et al., 2009). Based upon the results obtained, it might be necessary to investigate alternative instruments that could better elucidate the differences of interest. Going beyond specific instruments, other theoretical perspectives may be more equipped to explain political differences in discourse. Likewise, an atheoretical approach in which large quantities of data are collected from which theories are formulated may be best suited to this area of research.

In critiquing MFT, Schein and Gray (2015) proposed an alternative theory explaining moral choices and potential partisan differences: Dyadic Theory of Morality. Across many studies, researchers find that supposed differences in moral foundation can be more easily explained by a harm-focused morality where someone/thing is harmed by someone/thing else (Schein & Gray (2015)). Rather than liberals and conservatives having different conceptions of morality, dyadic morality argues that they simply identify harm differently. For example, in the case of abortion, liberals tend to identify the primary harm to the mother leading to a pro-choice position whereas conservatives tend to identify the primary harm to the fetus leading to a pro-life position. While as of yet, no linguistic measurement of dyadic morality exists; such an approach may be able to better identify the systematic differences between liberals and conservatives.

Turning to atheoretical approaches, future studies aiming to uncover political lean in discourse may benefit from a bottom-up, data-driven approach. Researchers might derive substantive insights into political lean through gathering data after which they may formulate theories that explain systematic observations obtained from that data. The availability of text corpora along with methods for extracting large amounts of text from the internet (as was demonstrated in the current study) potentially make this a feasible option. Likewise, there are several approaches to analyzing such data, including linear models (like multilevel models) and network-style models such as latent semantic analysis (Landauer, 1998). Owing to the wealth of representative data as well as the sophistication of current analytic tools, there is high potential for the explanatory power of new theories involving political discourse.

Never before has political discourse represented such fertile ground for psychological research. The plethora of options for news sources has created not only an abundance of choice but also vast quantities of text data. Along with this recent increase in the amount of text information, there is now an obligation on the part of researchers to devise proper

methods for analyzing that text. Solid methodologies must be constructed and periodically improved to keep pace with evolving technologies. Likewise, a strong theoretical foundation is paramount to making sense of the current and future media ecosystem. Therefore, it is incumbent upon social scientists to continue investigating the information consumed by millions of Americans every day so that insights into the nature and consequences of political discourse can be more completely understood.

References

- Aust, F., & Barth, M. (2017). *papaja: Create APA manuscripts with R Markdown*.
Retrieved from <https://github.com/crsh/papaja>
- Bentler, P. M. (1990). Comparative fit indexes in structural models. *Psychological Bulletin*, 107(2), 238–246. <https://doi.org/10.1037/0033-2909.107.2.238>
- Bentler, P. M., & Bonett, D. G. (1980). Significance tests and goodness of fit in the analysis of covariance structures. *Psychological Bulletin*, 88(3), 588–606.
<https://doi.org/10.1037/0033-2909.88.3.588>
- Bradley, M. M., & Lang, P. J. (1999). *Affective Norms for English Words (ANEW): Instruction Manual and Affective Ratings*. The Center for Research in Psychophysiology.
- Buuren, S. V., & Groothuis-Oudshoorn, K. (2011). **mice** : Multivariate Imputation by Chained Equations in R. *Journal of Statistical Software*, 45(3).
<https://doi.org/10.18637/jss.v045.i03>
- Byrne, B. M. (2001). Structural Equation Modeling With AMOS, EQS, and LISREL: Comparative Approaches to Testing for the Factorial Validity of a Measuring Instrument. *International Journal of Testing*, 1(1), 55–86.
https://doi.org/10.1207/S15327574IJT0101_4
- Clifford, S., & Jerit, J. (2013). How words do the work of politics: Moral foundations theory and the debate over stem cell research. *The Journal of Politics*, 75(3), 659–671.
<https://doi.org/10.1017/S0022381613000492>
- De Deyne, S., Navarro, D. J., Perfors, A., Brysbaert, M., & Storms, G. (2019). The “Small World of Words” English word association norms for over 12,000 cue words. *Behavior Research Methods*, 51(3), 987–1006. <https://doi.org/10.3758/s13428-018-1115-7>
- Federico, C. M., Weber, C. R., Ergun, D., & Hunt, C. (2013). Mapping the Connections between Politics and Morality: The Multiple Sociopolitical Orientations Involved in Moral Intuition. *Political Psychology*, 34(4), 589–610.
<https://doi.org/10.1111/pops.12006>

- Feinerer, I., & Hornik, K. (2017). *Text mining package*. Retrieved from <http://tm.r-forge.r-project.org/>
- Feinerer, I., Hornik, K., & Meyer, D. (2008). Text Mining Infrastructure in *R*. *Journal of Statistical Software*, 25(5). <https://doi.org/10.18637/jss.v025.i05>
- Frimer, J. A. (2020). Do liberals and conservatives use different moral languages? Two replications and six extensions of Graham, Haidt, and Nosek's (2009) moral text analysis. *Journal of Research in Personality*, 84, 103906. <https://doi.org/10.1016/j.jrp.2019.103906>
- Frimer, J. A., Tell, C. E., & Haidt, J. (2015). Liberals condemn sacrilege too: The harmless desecration of Cerro Torre. *Social Psychological and Personality Science*, 6(8), 878–886. <https://doi.org/10.1177/1948550615597974>
- Garten, J., Boghrati, R., Hoover, J., Johnson, K. M., & Dehghani, M. (2016). Morality between the lines: Detecting moral sentiment in text. *Proceedings of IJCAI 2016 workshop on Computational Modeling of Attitudes*.
- Gelman, A. (2006). Multilevel (hierarchical) modeling: What it can and cannot do. *Technometrics*, 48(3), 432–435. <https://doi.org/10.1198/004017005000000661>
- Graham, J., Haidt, J., & Nosek, B. A. (2009). Liberals and conservatives rely on different sets of moral foundations. *Journal of Personality and Social Psychology*, 96(5), 1029–1046. <https://doi.org/10.1037/a0015141>
- Graham, J., Nosek, B. A., & Haidt, J. (2012). The Moral Stereotypes of Liberals and Conservatives: Exaggeration of Differences across the Political Spectrum. *PLoS ONE*, 7(12), e50092. <https://doi.org/10.1371/journal.pone.0050092>
- Graham, J., Nosek, B. A., Haidt, J., Iyer, R., Koleva, S., & Ditto, P. H. (2011a). Mapping the moral domain. *Journal of Personality and Social Psychology*, 101(2), 366–385. <https://doi.org/10.1037/a0021847>
- Graham, J., Nosek, B. A., Haidt, J., Iyer, R., Koleva, S., & Ditto, P. H. (2011b). Mapping the moral domain. *Journal of Personality and Social Psychology*, 101(2), 366–385. <https://doi.org/10.1037/a0021847>

- Gray, K., & Keeney, J. E. (2015). Disconfirming moral foundations theory on its own terms: Reply to Graham (2015). *Social Psychological and Personality Science*, 6(8), 874–877. <https://doi.org/10.1177/1948550615592243>
- Gray, K., Schein, C., & Ward, A. F. (2014). The myth of harmless wrongs in moral cognition: Automatic dyadic completion from sin to suffering. *Journal of Experimental Psychology: General*, 143(4), 1600–1615. <https://doi.org/10.1037/a0036149>
- Haidt, J., & Graham, J. (2007). When morality opposes justice: Conservatives have moral intuitions that Liberals may not recognize. *Social Justice Research*, 20(1), 98–116. <https://doi.org/10.1007/s11211-007-0034-z>
- Haidt, J., & Joseph, C. (2004). Intuitive ethics: How innately prepared intuitions generate culturally variable virtues. *Daedalus*, 133(4), 55–66. Retrieved from <https://www.jstor.org/stable/20027945>
- Harrison, J., & Kim, J. Y. (2020). *RSelenium: R Bindings for 'Selenium WebDriver'*. Retrieved from <https://CRAN.R-project.org/package=RSelenium>
- Hopp, F. R., Fisher, J. T., Cornell, D., Huskey, R., & Weber, R. (2021). The extended moral foundations dictionary (eMFD): Development and applications of a crowd-sourced approach to extracting moral intuitions from text. *Behavior Research Methods*, 53(1), 232–246. <https://doi.org/10.3758/s13428-020-01433-0>
- Hu, L., & Bentler, P. M. (1999). Cutoff criteria for fit indexes in covariance structure analysis: Conventional criteria versus new alternatives. *Structural Equation Modeling: A Multidisciplinary Journal*, 6(1), 1–55. <https://doi.org/10.1080/10705519909540118>
- Hutto, C. (2021). *Welcome to VaderSentiment's documentation! — VaderSentiment 3.3.1 documentation*. Retrieved from <https://vadersentiment.readthedocs.io/en/latest/>
- Jöreskog, K. G. (1971). Simultaneous factor analysis in several populations. *Psychometrika*, 36(4), 409–426. <https://doi.org/10.1007/BF02291366>
- Kincaid, J. P., Fishburne, Jr., Robert P., R., Richard L., C., & Brad S. (1975). *Derivation of new readability formulas (automated readability index, fog count and Flesch reading ease*

- formula*) for Navy enlisted personnel. <https://doi.org/10.21236/ADA006655>
- Kohlberg, L., & Hersh, R. H. (1977). Moral development: A review of the theory. *Theory Into Practice*, 16(2), 53–59. <https://doi.org/10.1080/00405847709542675>
- Landauer, T. K. (1998). Learning and Representing Verbal Meaning. *Current Directions in Psychological Science*, 7(5), 161–164. <https://doi.org/10.1111/1467-8721.ep10836862>
- Leveau, N., Jhean-Larose, S., Denhière, G., & Nguyen, B.-L. (2012). Validating an interlingual metanorm for emotional analysis of texts. *Behavior Research Methods*, 44(4), 1007–1014. <https://doi.org/10.3758/s13428-012-0208-y>
- Loria, S. (2020). *Tutorial: Quickstart — TextBlob 0.16.0 documentation*. Retrieved from <https://textblob.readthedocs.io/en/dev/quickstart.html>
- Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Efficient Estimation of Word Representations in Vector Space. *arXiv:1301.3781 [Cs]*. Retrieved from <http://arxiv.org/abs/1301.3781>
- Mitchell, A., Matsa, K. E., Gottfried, J., & Kiley, J. (2014). *Political polarization & media habits* / Pew Research Center. Retrieved from <http://www.journalism.org/2014/10/21/political-polarization-media-habits/>
- Nelson, D. L., McEvoy, C. L., & Schreiber, T. A. (2004). The University of South Florida free association, rhyme, and word fragment norms. *Behavior Research Methods, Instruments, & Computers*, 36(3), 402–407. <https://doi.org/10.3758/BF03195588>
- Pennebaker, J. W., Booth, R. J., & Frances, M. E. (2007). *Liwc2007: Linguistic inquiry and word count*. Austin, TX.
- Perry, P. O. (2024). *Corpus: Text corpus analysis*. Retrieved from <https://leslie-huang.github.io/r-corpus/>
- Pinheiro, J., Bates, D., Debroy, S., Sarkar, D., & Team, R. C. (2017). *nlme: Linear and nonlinear mixed effects models*. Retrieved from <https://cran.r-project.org/package=nlme>
- Sagi, E., & Dehghani, M. (2014). Measuring moral rhetoric in text. *Social Science Computer Review*, 32(2), 132–144. <https://doi.org/10.1177/0894439313506837>

- Schein, C., & Gray, K. (2015). The unifying moral dyad: Liberals and conservatives share the same harm-based moral template. *Personality and Social Psychology Bulletin*, 41(8), 1147–1163. <https://doi.org/10.1177/0146167215591501>
- Schmidt, M. F. H., Gonzalez-Cabrera, I., & Tomasello, M. (2017). Children's developing metaethical judgments. *Journal of Experimental Child Psychology*, 164, 163–177. <https://doi.org/10.1016/j.jecp.2017.07.008>
- Steiger, J. H. (1990). Structural model evaluation and modification: An interval estimation approach. *Multivariate Behavioral Research*, 25(2), 173–180. https://doi.org/10.1207/s15327906mbr2502_4
- Sterling, J., & Jost, J. T. (2018). Moral discourse in the twitterverse: Effects of ideology and political sophistication on language use among US citizens and members of congress. *Journal of Language and Politics*, 17(2), 195–221. <https://doi.org/10.1075/jlp.17034.ste>
- Suhler, C. L., Churchland, P., & Joseph, C. (2011). Can innate , modular "foundations" explain morality? Challenges for Haidt's Moral Foundations Theory. *Journal of Cognitive Neuroscience*, 23(9), 2103–2116. <https://doi.org/10.1162/jocn.2011.21637>
- Tabachnick, B. G., & Fidell, L. S. (2012). *Using multivariate statistics* (Sixth). Boston, MA: Pearson.
- Warriner, A. B., Kuperman, V., & Brysbaert, M. (2013). Norms of valence, arousal, and dominance for 13,915 English lemmas. *Behavior Research Methods*, 45(4), 1191–1207. <https://doi.org/10.3758/s13428-012-0314-x>
- Weber, C. R., & Federico, C. M. (2013). Moral Foundations and Heterogeneity in Ideological Preferences. *Political Psychology*, 34(1), 107–126. <https://doi.org/10.1111/j.1467-9221.2012.00922.x>
- Wickham, H. (2016). *Package 'rvest'*. Retrieved from <https://cran.r-project.org/package=rvest>
- Widaman, K. F. (1985). Hierarchically Nested Covariance Structure Models for Multitrait-Multimethod Data. *Applied Psychological Measurement*, 9(1), 1–26.

<https://doi.org/10.1177/014662168500900101>

Table 1

*Mean Percent Scores and MFQ Scores for
Foundation Areas*

Foundation Area	Dictionary/MFQ	M	SD
Authority	MFQ	11.52	3.06
Authority	Original	0.37	0.62
Authority	Complete	3.88	2.12
Authority	Reduced	0.98	1.06
Fair	MFQ	14.12	2.78
Fair	Original	1.05	0.95
Fair	Complete	4.94	2.37
Fair	Reduced	3.16	1.85
Harm	MFQ	13.85	2.79
Harm	Original	1.13	1.21
Harm	Complete	7.17	3.05
Harm	Reduced	3.97	2.24
Ingroup	MFQ	11.96	3.24
Ingroup	Original	0.97	1.40
Ingroup	Complete	5.12	2.63
Ingroup	Reduced	1.69	1.67
Purity	MFQ	11.69	3.61
Purity	Original	0.73	0.82
Purity	Complete	3.92	2.20
Purity	Reduced	1.94	1.46

Note. Mean percents for dictionaries, mean total subscores for the MFQ.

Table 2

MTMM Fit Indices for Model 1 MTMM

Model	Chi-Sq	df	RMSEA	SRMR	CFI	TLI	AIC
Original	204.40	71.00	0.03	0.04	0.98	0.97	21,575.55
Complete	239.09	71.00	0.04	0.04	0.97	0.96	24,868.17
Reduced	283.10	71.00	0.05	0.05	0.95	0.94	23,380.58

Table 3

MTMM Fit Estimates for Model 1 Loadings

Latent Variable	Measured Variable	Original	All	Reduced
Harm	1	0.61	0.61	0.61
Harm	2	0.70	0.69	0.69
Harm	3	0.38	0.37	0.36
Harm	MFD	-0.05	-0.03	-0.05
Fair	1	0.76	0.76	0.77
Fair	2	0.52	0.50	0.50
Fair	3	0.23	0.22	0.22
Fair	MFD	-0.01	0.10	0.05
Ingroup	1	0.63	0.63	0.63
Ingroup	2	0.73	0.72	0.73
Ingroup	3	0.82	0.81	0.81
Ingroup	MFD	-0.10	0.03	0.05
Authority	1	0.81	0.80	0.81
Authority	2	0.72	0.72	0.72
Authority	3	0.42	0.41	0.41
Authority	MFD	0.06	0.00	-0.04
Purity	1	0.71	0.71	0.71
Purity	2	0.67	0.67	0.67
Purity	3	0.64	0.64	0.64
Purity	MFD	0.06	-0.05	-0.11

Note. Estimates are completely standardized loadings from the model.

Table 4

Example of Weighted Coding for Harm Words

Source	Concept	Frequency	Total Words	Valence	Weighted Score
Breitbart	defend	4.00	341	0.97	1.14
NPR	protect	11.00	1550	1.39	0.98
NPR	fight	7.00	10740	-1.20	-0.08
NPR	violenc	11.00	1419	-1.85	-1.43
Breitbart	safe	4.00	1108	2.07	0.75
Breitbart	sympathet	1.00	1632	1.26	0.08
Breitbart	benefit	6.00	1774	1.42	0.48
NY Times	care	19.00	4422	2.02	0.87
NY Times	harm	4.00	1100	-2.47	-0.90
Breitbart	damag	4.00	1624	-1.63	-0.40
Breitbart	preserv	2.00	254	1.27	1.00
NY Times	destroy	5.00	1494	-1.88	-0.63
NY Times	spurn	1.00	1508	NA	NA
Breitbart	cruel	2.00	574	-1.83	-0.64
Breitbart	violent	6.00	1211	-2.20	-1.09

Note. Concepts were stemmed to match datasets.

Table 5

Study 2 Results - Multilevel Model

Model	Foundation	M_C	SD_C	M_L	SD_L	t	p	d	ICC
Valence	Harm/Care	-0.49	1.26	-0.44	1.26	0.50	.668	-0.04	.001
Valence	Fairness/Reciprocity	0.47	0.64	0.44	0.64	-1.23	.344	0.05	< .001
Valence	Ingroup/Loyalty	0.39	0.49	0.38	0.49	-0.58	.620	0.02	< .001
Valence	Authority/Respect	0.09	0.49	0.07	0.49	-1.25	.338	0.05	< .001
Valence	Purity/Sanctity	0.19	0.53	0.13	0.53	-2.77	.109	0.10	.001
Original	Harm/Care	0.55	0.60	0.50	0.60	-1.84	.207	0.07	< .001
Original	Fairness/Reciprocity	0.46	0.49	0.41	0.49	-2.57	.124	0.10	.003
Original	Ingroup/Loyalty	0.63	0.64	0.68	0.64	1.83	.208	-0.07	.002
Original	Authority/Respect	0.67	0.70	0.60	0.70	-2.28	.150	0.09	< .001
Original	Purity/Sanctity	0.17	0.32	0.11	0.32	-3.78	.064	0.20	.009

Note. For mean and standard deviation values, 'C' and 'L' refer to 'conservative' and 'liberal,' respectively

Table 6

Kavanaugh Results - Multilevel Model

Model	Foundation	M_C	SD_C	M_L	SD_L	t	p	d	ICC
Valence	Harm/Care	-0.35	0.69	-0.32	0.69	0.58	.575	-0.04	.013
Valence	Fairness/Reciprocity	0.91	0.74	0.60	0.74	-2.68	.028	0.42	.095
Valence	Ingroup/Loyalty	0.24	0.31	0.25	0.31	0.31	.764	-0.03	.004
Valence	Authority/Respect	0.04	0.32	0.06	0.32	1.08	.313	-0.06	.003
Valence	Purity/Sanctity	0.36	0.51	0.16	0.51	-1.75	.118	0.47	.133
Original	Harm/Care	0.33	0.36	0.35	0.36	0.60	.567	-0.04	.010
Original	Fairness/Reciprocity	0.77	0.57	0.54	0.57	-2.84	.022	0.40	.142
Original	Ingroup/Loyalty	0.34	0.37	0.37	0.37	0.35	.733	-0.06	.010
Original	Authority/Respect	0.43	0.49	0.46	0.49	1.20	.266	-0.04	.004
Original	Purity/Sanctity	0.28	0.31	0.15	0.31	-2.44	.040	0.49	.126

Note. For mean and standard deviation values, 'C' and 'L' refer to 'conservative' and 'liberal,' respectively

Table 7

Government Shutdown Results - Multilevel Model

Model	Foundation	M_C	SD_C	M_L	SD_L	t	p	d	ICC
Valence	Harm/Care	-0.13	0.70	-0.15	0.70	-0.55	.594	0.03	.008
Valence	Fairness/Reciprocity	0.79	0.57	0.40	0.57	-1.57	.154	0.74	.139
Valence	Ingroup/Loyalty	0.35	0.39	0.37	0.39	0.53	.611	-0.05	.012
Valence	Authority/Respect	0.03	0.33	0.03	0.33	0.67	.524	-0.01	.004
Valence	Purity/Sanctity	0.48	0.49	0.08	0.49	-1.38	.206	0.85	.155
Original	Harm/Care	0.34	0.38	0.36	0.38	1.24	.251	-0.04	< .001
Original	Fairness/Reciprocity	0.68	0.44	0.38	0.44	-2.35	.047	0.75	.148
Original	Ingroup/Loyalty	0.51	0.50	0.54	0.50	0.23	.825	-0.06	.015
Original	Authority/Respect	0.42	0.44	0.43	0.44	0.58	.578	-0.02	.002
Original	Purity/Sanctity	0.34	0.30	0.13	0.30	-1.38	.206	0.85	.155

Note. For mean and standard deviation values, 'C' and 'L' refer to 'conservative' and 'liberal,' respectively

Appendix A

Study 2: Descriptive Statistics

We calculated descriptive statistics for each news source in order to understand any and all fundamental linguistic differences in the sources' use of English. Statistics calculated included average z -scored valence of the unique words per article, number of articles per source, total number of words per source, average number of tokens (words) per article in each source, average number of types (unique words) per article in each source, and mean readability level per source. Readability statistics were calculated using the Flesch-Kincaid Grade Level Readability formula (Kincaid, Fishburne, Robert P., Richard L., & Brad S., 1975). Readability is calculated using a formula where the total number of syllables, words, and sentences in a given passage are determinants of its difficulty. The obtained value is intended to match up with the U.S. grade level at which one should be able to comfortably read the passage (Kincaid et al., 1975). For example, a text with a readability score of 11 should be easily read by a U.S. high school junior.

As seen in Table A1, the sources are similar in some aspects yet different in others. Valence appears to be slightly positive across all sources. The large standard deviations seem to indicate little to no presence of a difference in valence across sources. *NYT* published the greatest number of articles as well as total words. *Breitbart* featured the lowest number of articles, and *NPR* the lowest number of total words from all articles. Per individual article, however, *Breitbart* appears to feature the highest average number of words as well as unique words. Once again the standard deviations call into question any apparent differences between sources. Finally, *Fox News* articles had the lowest reading grade level on average, while *NYT* had the highest. This result might be attributable to the greater number of tokens in the average *NYT* article compared to *Fox News*. The standard deviations for readability suggest the presence of a diverse array of articles for each source, ranging from low to high reading level. Large standard deviations suggest the sources feature a lot of

overlap between them in their representation of scores.

Table A1

Study 2 - Descriptive Statistics by Source

Source	M_V	SD_V	$N_{Article}$	N_{Words}	M_T	SD_T	M_{Ty}	SD_{Ty}	M_{FK}	SD_{FK}
NPR	0.28	0.23	695	302977	435.94	642.63	191.96	192.28	14.00	3.93
New York Times	0.30	0.13	406	452579	1114.73	511.86	454.27	154.58	16.44	3.36
Breitbart	0.29	0.18	1437	722022	502.45	347.90	243.35	120.75	18.56	7.90
Fox News	0.29	0.17	503	296779	590.02	528.60	283.56	189.00	17.25	7.21

Note. Readability statistics were calculated using the Flesch-Kincaid Grade Level readability formula. V = Valence, T = Tokens or total words, Ty = Types or unique words, FK = Flesch-Kincaid

Appendix B

Study 3: Descriptive Statistics

We calculated descriptive statistics for each news source per topic in order to reveal the presence (if any) of linguistic differences in the sources' use of language. As in Study 1, statistics calculated include z -scored valence, number of articles per source, total words per source, mean tokens per article in each source, mean types per article in each source, and mean readability level (using the Flesch-Kincaid Grade Level Readability formula) per source (Kincaid et al., 1975).

Table B1 displays the descriptive statistics for sources' writing on the Kavanaugh hearing confirmation hearing. The sources were similar in most basic linguistic aspects, except for number of articles. For example, *Sean Hannity* appears to have published only 27 articles while *Breitbart* published 757 articles on this topic. Valence was found to be slightly positive across all sources. *Fox News* produced the most total words with the most tokens on average. This is likely due to the fact *Fox News* transcribes many of their videos and publishes them in article form. *Politico* featured the highest number of types on average. *Rush Limbaugh* featured the lowest readability score on average by grade level while *Slate* featured the highest grade-level readability score. The large standard deviations for these statistics, however, preclude conclusions regarding differences in the sources' use of language, as there is likely a lot of overlap between sources' use of language.

Table B2 displays descriptive statistics for articles about the partial government shutdown. Like the Kavanaugh hearing, the sources were similar in average valence (slightly positive). Once again, there was variation in the number of articles published by each source on this topic. *Sean Hannity*, *Rush Limbaugh*, and *The Blaze* published fewer than 100 articles while *Fox News* published 1,013 articles. *Fox News* again featured the most total words and mean tokens, but this is likely due to the presence of a large amount of video transcriptions that the organization published as articles. *Politico* had the most types on

Table B1

Kavanaugh - Descriptive Statistics by Source

Source	M_V	SD_V	$N_{Article}$	N_{Words}	M_T	SD_T	M_{Ty}	SD_{Ty}	M_{FK}	SD_{FK}
Huffington Post	0.27	0.12	552	359046	650.45	462.21	283.07	129.23	10.68	1.89
NPR	0.22	0.25	366	108605	296.73	499.10	128.12	172.51	12.55	3.67
New York Times	0.33	0.13	653	723569	1108.07	570.31	461.26	174.20	9.50	1.79
Politico	0.32	0.10	689	1069292	1551.95	1045.47	614.66	358.68	12.09	2.66
Slate	0.29	0.12	272	229896	845.21	612.22	332.42	168.17	12.29	2.48
The Blaze	0.25	0.13	277	128097	462.44	149.66	210.52	53.14	10.73	1.85
Breitbart	0.30	0.13	757	375848	496.50	609.91	230.03	153.59	11.00	2.32
Fox News	0.29	0.11	646	1304048	2018.65	2709.77	534.46	404.20	9.91	1.98
Sean Hannity	0.31	0.11	27	5926	219.48	145.98	121.04	57.26	11.76	2.40
Rush Limbaugh	0.38	0.13	172	267067	1552.72	1258.51	419.00	225.80	9.17	8.86

Note. Readability statistics were calculated using the Flesch-Kincaid Grade Level readability formula. V = Valence, T = Tokens or total words, Ty = Types or unique words, FK = Flesch-Kincaid

average. For this topic, *Fox News* featured the lowest reading grade level while *Slate* featured the highest reading grade level. For each statistic, the excessively high standard deviations render any assertions regarding linguistic differences inconclusive on a descriptive level due to the aforementioned overlap in sources' language use.

Table B2

Government Shutdown - Descriptive Statistics by Source

Source	M_V	SD_V	$N_{Article}$	N_{Words}	M_T	SD_T	M_{Ty}	SD_{Ty}	M_{FK}	SD_{FK}
Huffington Post	0.30	0.13	432	242220	560.69	425.52	258.27	130.54	10.67	1.85
NPR	0.23	0.22	434	151114	348.19	457.34	155.53	171.30	11.20	3.16
New York Times	0.33	0.13	752	836811	1112.78	532.72	472.95	177.57	10.09	1.74
Politico	0.29	0.10	222	349499	1574.32	875.63	638.32	313.34	11.30	1.26
Slate	0.30	0.12	117	85254	728.67	464.34	301.71	138.74	11.86	2.54
The Blaze	0.26	0.13	98	38983	397.79	102.10	185.66	41.93	10.70	1.85
Breitbart	0.34	0.15	309	102886	332.96	238.70	168.31	81.59	10.70	2.04
Fox News	0.35	0.13	1013	2799217	2763.29	3405.44	637.28	493.98	9.61	2.11
Sean Hannity	0.29	0.14	63	10311	163.67	33.59	95.13	16.59	13.36	4.81
Rush Limbaugh	0.37	0.12	78	152630	1956.79	1544.14	482.79	252.79	9.77	7.64

Note. Readability statistics were calculated using the Flesch-Kincaid Grade Level readability formula. V = Valence, T = Tokens or total words, Ty = Types or unique words, FK = Flesch-Kincaid